

Automated Classification of Retinal Fundus Images for Diabetic Detection Using Machine Learning and Hyperparameter Optimization

Submitted by

Name	Roll Number
Aditya Jyosyula	3122225001006
K. Harish	3122225001036
Jagrath Mithun S	3122225001044
Janeshvar Sivakumar	3122225001047

Department of Computer Science and Engineering

SSN College of Engineering

April 8, 2025

Abstract

This project focuses on the classification of retinal fundus images for the detection of diabetes using machine learning models. Dimensionality reduction techniques and hyperparameter optimization were applied to enhance model performance. The project uses the Retina Blood Vessel dataset, specifically focusing on diabetic cases. Results show that optimized models significantly outperform their non-optimized counterparts.

1. Introduction

Diabetes is one of the leading causes of preventable blindness worldwide. Retinal fundus images provide critical information for the early detection of diabetic retinopathy. Manual inspection is time-consuming and requires expertise. Automating this task using machine learning can improve diagnostic efficiency and accessibility.

1.1. Problem Statement

To develop an optimized machine learning model capable of accurately classifying retinal fundus images as diabetic or non-diabetic using dimensionality reduction and hyperparameter tuning techniques.

2. Literature Survey

Several studies have explored diabetic retinopathy detection using deep learning and traditional machine learning methods. CNNs have shown strong performance in image-based tasks. Recent works emphasize the importance of preprocessing, feature extraction, and parameter tuning for improved model accuracy.

Miri et al. (2017) mention that inhomogeneous contrast and illumination, caused by imaging processes and biological characteristics, are significant challenges in retinal image analysis. They point out that color variations within and across images arise from factors like hemoglobin oxygen saturation, aging, cataract development, flash intensity and spectrum, camera distortions, flash artifacts, and focus.

Zhang et al. (2023) focuses on a particular preprocessing method called adaptive contrast enhancement (ACE) applied to the RITE dataset. They explain that ACE addresses the issues of uneven lighting, low overall illumination, and low contrast between blood vessels and the background, which are common in retinal fundus images and can negatively impact classification accuracy.

Ajaz et al. specifically focus on the relationship between retinal vessel geometrical features and the incidence and progression of DME. They found that Average Branching Angle (ABA) was the only parameter that exhibited a monotonic increase with disease severity.

Guo et al. investigate the association between retinal information and CVD in patients with type 2 diabetes, independent of traditional cardiovascular risk factors.

3. Proposed System

The proposed system includes a pipeline consisting of data preprocessing, feature engineering using PCA, model training, hyperparameter tuning via Grid Search and Random Search, and performance evaluation.

3.1. System Architecture Diagram




Figure 1: System Architecture Diagram

4. Implementation and Experiments

4.1. Development Environment

- Python 3.10
- Libraries: scikit-learn, pandas, matplotlib, seaborn, numpy, OpenCV
- Jupyter Notebook / VS Code
- Cloud Deployment: GCP (Cloud Run + Firebase Hosting)

4.2. Dataset Description

Source: Retina Blood Vessel Dataset on Kaggle

Focus: Only diabetes-labeled retinal images were used. The dataset contains high-resolution fundus images.

5. Implementation and Experiments

5.1. Development Environment

- Python 3.10
- Libraries: scikit-learn, pandas, matplotlib, seaborn, numpy, OpenCV
- Jupyter Notebook / VS Code
- Cloud Deployment: GCP (Cloud Run + Firebase Hosting)

5.2. Dataset Description

Source: Retina Blood Vessel Dataset on Kaggle

Focus: Only diabetes-labeled retinal images were used. The dataset contains high-resolution fundus images.

5.3. Implementation Steps

1. Data Collection:

Images were extracted from the Kaggle dataset. All diabetic-labeled images were filtered and selected for this binary classification task.

2. Preprocessing:

Each image underwent several enhancement techniques:

- Green channel extraction (retinal vessels are clearer in this channel)
- Contrast Limited Adaptive Histogram Equalization (CLAHE)
- Image resizing to 224×224
- Normalization to $[0, 1]$ scale

3. Feature Engineering:

Flattened image pixels formed high-dimensional feature vectors (over 50,000 features per image). PCA was applied to reduce dimensionality while retaining essential variance. The number of components was chosen based on explained variance threshold (e.g., 95%).

4. Model Training:

We trained the following ML models:

- Logistic Regression: A linear baseline model
- SVM: Effective in high-dimensional space
- Random Forest: Captures non-linear relations with ensemble learning

5. Hyperparameter Optimization:

We used:

- Grid Search: Exhaustive parameter tuning
- Randomized Search: Faster alternative with good results

Both used cross-validation and prioritized F1-score.


6. Model Evaluation:

We evaluated the models using:

- Accuracy, Precision, Recall, F1-score
- Confusion Matrix
- Cross-validation statistics

Classification Report for Simple CNN:				
	precision	recall	f1-score	support
0	0.58	0.66	0.62	140
1	0.61	0.51	0.56	140
accuracy				0.59
macro avg	0.59	0.59	0.59	280
weighted avg	0.59	0.59	0.59	280

Accuracy: 0.5892857142857143




```

Classification Report for EfficientNetB0:
precision    recall   f1-score   support
          0       0.54      0.87      0.66      140
          1       0.66      0.25      0.36      140

accuracy                           0.56      280
macro avg       0.60      0.56      0.51      280
weighted avg    0.60      0.56      0.51      280

Accuracy: 0.5607142857142857

```




```

Classification Report for ResNet50:
precision    recall   f1-score   support
          0       0.65      0.14      0.23      140
          1       0.52      0.92      0.66      140

accuracy                           0.53      280
macro avg       0.58      0.53      0.45      280
weighted avg    0.58      0.53      0.45      280

Accuracy: 0.5321428571428571

```




```

Classification Report for VGG16:
precision    recall   f1-score   support
          0       0.56      0.59      0.58      140
          1       0.57      0.54      0.56      140

accuracy                           0.57      280
macro avg       0.57      0.57      0.57      280
weighted avg    0.57      0.57      0.57      280

Accuracy: 0.5678571428571428

```



7. Model Ensembling:

We combined the SVM and Random Forest models for better generalization. This boosted the performance significantly over individual models.

Evaluation with PCA Features:					
• Logistic Regression Accuracy: 0.5857					
	precision	recall	f1-score	support	
0	0.59	0.56	0.57	140	
1	0.58	0.61	0.60	140	
accuracy			0.59	280	
macro avg	0.59	0.59	0.59	280	
weighted avg	0.59	0.59	0.59	280	
• SVM Accuracy: 0.5821					
	precision	recall	f1-score	support	
0	0.59	0.52	0.56	140	
1	0.57	0.64	0.61	140	
accuracy			0.58	280	
macro avg	0.58	0.58	0.58	280	
weighted avg	0.58	0.58	0.58	280	
• Random Forest Accuracy: 0.5750					
	precision	recall	f1-score	support	
0	0.57	0.61	0.59	140	
1	0.58	0.54	0.56	140	
accuracy			0.57	280	
macro avg	0.58	0.57	0.57	280	
weighted avg	0.58	0.57	0.57	280	

8. Deployment:

The final model was saved using `joblib`, containerized using Docker, and deployed to Google Cloud Platform (GCP) via Cloud Run. Firebase Hosting was used to serve a frontend UI for public access.

- **Deployed API:** <https://diabetes-api-356779724219.us-central1.run.app/>
- **GitHub Repository:** <https://github.com/juicjaane/ML-project-retina.git>

6. Results and Discussions




Figure 2: Performance Metrics of Optimized Models after ensemble

⌚ Evaluation with PCA Features:				
• Logistic Regression Accuracy: 0.5857				
precision	recall	f1-score	support	
0 0.59	0.56	0.57	140	
1 0.58	0.61	0.60	140	
accuracy		0.59	280	
macro avg	0.59	0.59	0.59	280
weighted avg	0.59	0.59	0.59	280
• SVM Accuracy: 0.5821				
precision	recall	f1-score	support	
0 0.59	0.52	0.56	140	
1 0.57	0.64	0.61	140	
accuracy		0.58	280	
macro avg	0.58	0.58	0.58	280
weighted avg	0.58	0.58	0.58	280
• Random Forest Accuracy: 0.5750				
precision	recall	f1-score	support	
0 0.57	0.61	0.59	140	
1 0.58	0.54	0.56	140	
accuracy		0.57	280	
macro avg	0.58	0.57	0.57	280
weighted avg	0.58	0.57	0.57	280

Figure 3: Classification Report and Confusion Matrix before ensemble

6.1. Human, Societal, Ethical, and Sustainable Impact

- Enables early diagnosis and reduces clinician burden.
- Equitable access in low-resource settings.
- Ethical AI practices focused on fairness and transparency.
- Supports SDG Goal 3: Good Health and Well-being.

7. Conclusion and Future Work

This project demonstrated the efficacy of combining dimensionality reduction and hyper-parameter tuning for classifying diabetic retinal images. Future work involves using deep

learning models like CNNs, integrating explainable AI tools (Grad-CAM), and expanding the system to detect severity levels of diabetic retinopathy.

8. References

1. Abdallah Wagih, Retina Blood Vessel Dataset, Kaggle, <https://www.kaggle.com/datasets/abdallahwagih/retina-blood-vessel>
2. Miri M, Amini Z, RabbaniH, Kafieh R. A comprehensive study of retinal vessel classification methods in fundus images. J Med Sign Sens 2017;7:59-70.
3. Aqsa Ajaz, Himeesh Kumar, Behzad Aliahmad, Dinesh K. Kumar, The relationship between retinal vessel geometrical changes to incidence and progression of Diabetic Macular Edema, Informatics in Medicine Unlocked, Volume 16, 2019.
4. Guo, V., Chan, J., Chung, H. et al. Retinal Information is Independently Associated with Cardiovascular Disease in Patients with Type 2 diabetes. Sci Rep 6, 19053 (2016).
5. Zhang, J. et al. End-to-End Automatic Classification of Retinal Vessel Based on GANs with Improved U-Net. Diagnostics 2023, 13, 1148.
6. WHO Diabetic Retinopathy Factsheet
7. Papers With Code - Diabetic Retinopathy Detection Benchmarks
8. Scikit-learn Documentation
9. GCP Cloud Run and Firebase Hosting Docs

A. Complete Code Listing

```
1 # ===== Cell 1 =====
2
3
4 # ===== Cell 2 =====
5 !pip install umap-learn
6
7 # ===== Cell 3 =====
8 import os
9 import cv2
10 import numpy as np
11 import pandas as pd
12 import matplotlib.pyplot as plt
13 from sklearn.model_selection import train_test_split
14 from sklearn.decomposition import PCA
15 import umap
16 from sklearn.metrics import accuracy_score, precision_score,
   recall_score, f1_score
17 from sklearn.ensemble import VotingClassifier
18 import tensorflow as tf
19 from tensorflow.keras import layers, models
20 from tensorflow.keras.utils import to_categorical
```

```

21 from tensorflow.keras.applications import VGG16, ResNet50,
22     EfficientNetB0
23 from tensorflow.keras.callbacks import EarlyStopping
24
25 # ===== Cell 4 =====
26 import os
27 import cv2
28 import numpy as np
29 import matplotlib.pyplot as plt
30
31 # Function for CLAHE application
32 def apply_clahe(image):
33     clahe = cv2.createCLAHE(clipLimit=2.0, tileGridSize=(8, 8))
34     return clahe.apply(image)
35
36 # Matched filter function
37 def matched_filter(image, kernel_sizes=[9, 15, 21], sigma=2, num_angles
38 =24, wavelengths=[10, 15], gammas=[0.5, 1]):
39     filtered_images = []
40     for kernel_size in kernel_sizes:
41         for angle in range(num_angles):
42             theta = np.deg2rad(angle * 180 / num_angles)
43             for wavelength in wavelengths:
44                 for gamma in gammas:
45                     kernel = cv2.getGaborKernel((kernel_size,
46                         kernel_size), sigma, theta, wavelength, gamma,
47                         0, ktype=cv2.CV_32F)
48                     filtered_image = cv2.filter2D(image, cv2.CV_32F,
49                         kernel)
50                     filtered_images.append(filtered_image)
51     combined_image = np.max(np.array(filtered_images), axis=0)
52     return combined_image
53
54 # ===== Cell 5 =====
55 def preprocess_image(image_path, target_size=(224, 224)):
56     image = cv2.imread(image_path)
57     if image is None:
58         print(f"Error: Could not load image {image_path}")
59         return None
60
61     # Step 1: Extract the green channel
62     green_channel = image[:, :, 1]
63
64     # Step 2: Apply Laplacian sharpening
65     log_image = cv2.Laplacian(green_channel, cv2.CV_64F)
66     log_image = np.uint8(255 * (log_image - np.min(log_image)) / (np.
67         max(log_image) - np.min(log_image)))
68
69     # Step 3: Apply CLAHE
70     clahe = apply_clahe(log_image)
71
72     # Step 4: Apply matched filter
73     matched = matched_filter(clahe, num_angles=10, wavelengths=[10],
74         gammas=[0.5, 1, 2])
75     matched = cv2.normalize(matched, None, 0, 255, cv2.NORM_MINMAX).
76         astype(np.uint8)

```

```

71
72     # Step 5: Apply Gaussian blur for unsharp masking
73     blurred = cv2.GaussianBlur(matched, (7, 7), 10.0)
74     unsharp_image = cv2.addWeighted(matched, 1.5, blurred, -0.5, 0)
75
76     # Step 6: Apply Non-Local Means Filtering
77     non_local_means_filtered = cv2.fastNlMeansDenoising(matched, None,
78                 h=3, templateWindowSize=3, searchWindowSize=50)
79
80     # Step 7: Apply morphological closing
81     kernel = np.ones((3, 3), np.uint8)
82     closed_image = cv2.morphologyEx(non_local_means_filtered, cv2.
83         MORPH_CLOSE, kernel)
84     closed_image = cv2.bitwise_not(closed_image)
85
86     # Step 8: Apply adaptive thresholding
87     adaptive_thresh_mean = cv2.adaptiveThreshold(
88         closed_image,
89         255,
90         cv2.ADAPTIVE_THRESH_MEAN_C,
91         cv2.THRESH_BINARY,
92         3,
93         2
94     )
95
96     # Step 9: Invert the image for connected component analysis
97     inverted = cv2.bitwise_not(adaptive_thresh_mean)
98     num_labels, labels_im = cv2.connectedComponents(inverted)
99     output_image = np.zeros_like(inverted)
100    min_size = 15
101    for label in range(1, num_labels):
102        component_size = np.sum(labels_im == label)
103        if component_size >= min_size:
104            output_image[labels_im == label] = 255
105
106    return output_image
107
108 # ===== Cell 6 =====
109 import os
110 import cv2
111 import numpy as np
112 import pandas as pd
113 from sklearn.model_selection import train_test_split
114 from sklearn.decomposition import PCA
115 from sklearn.metrics import classification_report, accuracy_score
116 from sklearn.ensemble import VotingClassifier
117 import tensorflow as tf
118 from tensorflow.keras import layers, models
119 from tensorflow.keras.utils import to_categorical
120 from tensorflow.keras.applications import VGG16, ResNet50,
121     EfficientNetB0
122 from tensorflow.keras.applications.vgg16 import preprocess_input as
123     vgg_preprocess
124 from tensorflow.keras.applications.resnet50 import preprocess_input as
125     resnet_preprocess
126 from tensorflow.keras.applications.efficientnet import preprocess_input
127     as efficientnet_preprocess
128 from sklearn.linear_model import LogisticRegression

```

```

123 from sklearn.svm import SVC
124 import umap
125
126 from sklearn.metrics import accuracy_score, precision_score,
127     recall_score, f1_score
128
129 # Load binary labeled data
130 def load_binary_data(base_path, target_size=(224, 224)):
131     file_paths, labels = [], []
132     label_map = {'N': 0, 'D': 1} # 0: Normal, 1: Diabetic
133
134     for label, folder in label_map.items():
135         folder_path = os.path.join(base_path, 'normal' if label == 'N'
136             else 'diabetes')
137         for fname in os.listdir(folder_path):
138             if fname.lower().endswith('.jpg', '.jpeg', '.png')):
139                 file_paths.append(os.path.join(folder_path, fname))
140                 labels.append(label_map[label])
141
142     # Preprocess images with progress updates every 10 images
143     images, valid_labels = [], []
144     for i, (path, label) in enumerate(zip(file_paths, labels)):
145         img = preprocess_image(path, target_size)
146         if img is not None:
147             images.append(img)
148             valid_labels.append(label)
149         if (i + 1) % 10 == 0:
150             print(f"    Processed {i + 1} images...")
151
152     X = np.expand_dims(np.array(images), axis=-1)
153     y = to_categorical(valid_labels, 2)
154     return train_test_split(X, y, test_size=0.2, random_state=42,
155                             stratify=y)
156
157 # Load data
158 base_dir = r"/kaggle/input/classificationmlprojext/datasets"
159 X_train, X_test, y_train, y_test = load_binary_data(base_dir)
160 #
161
162 # ===== Cell 7 =====
163 print(X_train)
164
165 # ===== Cell 8 =====
166 import tensorflow as tf
167 from tensorflow.keras import layers, models, optimizers
168 from tensorflow.keras.callbacks import EarlyStopping
169
170 # Build a simple CNN model
171 def build_simple_cnn(input_shape=(224, 224, 1), num_classes=2):
172     model = models.Sequential([
173         layers.Conv2D(32, (3, 3), activation='relu', input_shape=
174             input_shape),
175         layers.MaxPooling2D((2, 2)),
176         layers.Conv2D(64, (3, 3), activation='relu'),

```

```

177     layers.MaxPooling2D((2, 2)),
178
179     layers.Conv2D(128, (3, 3), activation='relu'),
180     layers.MaxPooling2D((2, 2)),
181
182     layers.Flatten(),
183     layers.Dense(128, activation='relu'),
184     layers.Dropout(0.5),
185     layers.Dense(num_classes, activation='softmax')
186 )
187
188 model.compile(
189     optimizer='adam',
190     loss='categorical_crossentropy',
191     metrics=['accuracy']
192 )
193 return model
194
195 # Instantiate and train
196 simple_cnn = build_simple_cnn()
197
198 # Early stopping
199 early_stop = EarlyStopping(monitor='val_loss', patience=5,
200                             restore_best_weights=True)
201
202 history_cnn = simple_cnn.fit(
203     X_train, y_train,
204     validation_data=(X_test, y_test),
205     epochs=30,
206     batch_size=32,
207     callbacks=[early_stop],
208     verbose=1
209 )
210
211 # Evaluate
212 cnn_preds = simple_cnn.predict(X_test)
213 cnn_preds_classes = tf.argmax(cnn_preds, axis=1)
214 y_true = tf.argmax(y_test, axis=1)
215
216 print("      Classification Report for Simple CNN:\n")
217 from sklearn.metrics import classification_report, accuracy_score
218 print(classification_report(y_true, cnn_preds_classes))
219 print("      Accuracy:", accuracy_score(y_true, cnn_preds_classes))
220
221 # ===== Cell 9 =====
222 import matplotlib.pyplot as plt
223
224 # Plot training & validation accuracy and loss
225 def plot_history(history, model_name="Simple CNN"):
226     acc = history.history['accuracy']
227     val_acc = history.history['val_accuracy']
228     loss = history.history['loss']
229     val_loss = history.history['val_loss']
230     epochs = range(1, len(acc) + 1)
231
232     plt.figure(figsize=(14, 5))

```

```

234     plt.subplot(1, 2, 1)
235     plt.plot(epochs, acc, 'b-', label='Training Acc')
236     plt.plot(epochs, val_acc, 'g-', label='Validation Acc')
237     plt.title(f'{model_name} - Accuracy')
238     plt.xlabel('Epochs')
239     plt.ylabel('Accuracy')
240     plt.legend()
241
242     plt.subplot(1, 2, 2)
243     plt.plot(epochs, loss, 'b-', label='Training Loss')
244     plt.plot(epochs, val_loss, 'g-', label='Validation Loss')
245     plt.title(f'{model_name} - Loss')
246     plt.xlabel('Epochs')
247     plt.ylabel('Loss')
248     plt.legend()
249
250     plt.show()
251
252 plot_history(history_cnn)
253
254
255 # ===== Cell 10 =====
256 from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
257
258 cm = confusion_matrix(y_true, cnn_preds_classes)
259 disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=["Normal", "Diabetic"])
260 disp.plot(cmap='Blues')
261 plt.title("Confusion Matrix - Simple CNN")
262 plt.show()
263
264
265 # ===== Cell 11 =====
266 from sklearn.metrics import roc_curve, auc
267
268 fpr, tpr, _ = roc_curve(y_test[:, 1], cnn_preds[:, 1])
269 roc_auc = auc(fpr, tpr)
270
271 plt.figure(figsize=(6, 6))
272 plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area = %0.4f)' % roc_auc)
273 plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
274 plt.title('ROC Curve - Simple CNN')
275 plt.xlabel('False Positive Rate')
276 plt.ylabel('True Positive Rate')
277 plt.legend(loc="lower right")
278 plt.grid()
279 plt.show()
280
281
282 # ===== Cell 12 =====
283
284
285 # ===== Cell 13 =====
286 from tensorflow.keras.applications import VGG16
287 from tensorflow.keras.applications.vgg16 import preprocess_input as vgg_preprocess
288 from tensorflow.keras import Model, Input

```

```

289
290 # Convert grayscale to 3-channel
291 def expand_channels(X):
292     return np.repeat(X, 3, axis=-1)
293
294 X_train_vgg = expand_channels(X_train)
295 X_test_vgg = expand_channels(X_test)
296
297 # Apply VGG preprocessing
298 X_train_vgg = vgg_preprocess(X_train_vgg)
299 X_test_vgg = vgg_preprocess(X_test_vgg)
300
301 # Build VGG16 model with custom head
302 def build_vgg16_model(input_shape=(224, 224, 3), num_classes=2):
303     base_model = VGG16(weights='imagenet', include_top=False,
304                         input_shape=input_shape)
305     base_model.trainable = False # Freeze layers
306
307     x = layers.GlobalAveragePooling2D()(base_model.output)
308     x = layers.Dense(128, activation='relu')(x)
309     x = layers.Dropout(0.5)(x)
310     outputs = layers.Dense(num_classes, activation='softmax')(x)
311
312     model = Model(inputs=base_model.input, outputs=outputs)
313     model.compile(optimizer='adam', loss='categorical_crossentropy',
314                     metrics=['accuracy'])
315     return model
316
317 vgg_model = build_vgg16_model()
318
319 history_vgg = vgg_model.fit(
320     X_train_vgg, y_train,
321     validation_data=(X_test_vgg, y_test),
322     epochs=10,
323     batch_size=32,
324     callbacks=[early_stop],
325     verbose=1
326 )
327
328 # ===== Cell 14 =====
329 # Predict and evaluate
330 vgg_preds = vgg_model.predict(X_test_vgg)
331 vgg_preds_classes = tf.argmax(vgg_preds, axis=1)
332
333 print("      Classification Report for VGG16:")
334 print(classification_report(y_true, vgg_preds_classes))
335 print("      Accuracy:", accuracy_score(y_true, vgg_preds_classes))
336
337 # ===== Cell 15 =====
338 plot_history(history_vgg, model_name="VGG16")
339
340
341 # ===== Cell 16 =====
342 cm = confusion_matrix(y_true, vgg_preds_classes)
343 disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=["Normal", "Diabetic"])

```

```

344 disp.plot(cmap='Purples')
345 plt.title("Confusion Matrix - VGG16")
346 plt.show()
347
348
349 # ===== Cell 17 =====
350 fpr, tpr, _ = roc_curve(y_test[:, 1], vgg_preds[:, 1])
351 roc_auc = auc(fpr, tpr)
352
353 plt.figure(figsize=(6, 6))
354 plt.plot(fpr, tpr, color='darkgreen', lw=2, label='ROC curve (area = %0.4f)' % roc_auc)
355 plt.plot([0, 1], [0, 1], color='gray', lw=2, linestyle='--')
356 plt.title('ROC Curve - VGG16')
357 plt.xlabel('False Positive Rate')
358 plt.ylabel('True Positive Rate')
359 plt.legend(loc="lower right")
360 plt.grid()
361 plt.show()
362
363
364 # ===== Cell 18 =====
365 from tensorflow.keras.applications import ResNet50
366 from tensorflow.keras.applications.resnet50 import preprocess_input as resnet_preprocess
367
368 # Prepare input
369 X_train_resnet = expand_channels(X_train)
370 X_test_resnet = expand_channels(X_test)
371
372 X_train_resnet = resnet_preprocess(X_train_resnet)
373 X_test_resnet = resnet_preprocess(X_test_resnet)
374
375 # Build ResNet50 model
376 def build_resnet50_model(input_shape=(224, 224, 3), num_classes=2):
377     base_model = ResNet50(weights='imagenet', include_top=False,
378                           input_shape=input_shape)
379     base_model.trainable = False # Freeze all layers
380
381     x = layers.GlobalAveragePooling2D()(base_model.output)
382     x = layers.Dense(128, activation='relu')(x)
383     x = layers.Dropout(0.5)(x)
384     outputs = layers.Dense(num_classes, activation='softmax')(x)
385
386     model = Model(inputs=base_model.input, outputs=outputs)
387     model.compile(optimizer='adam', loss='categorical_crossentropy',
388                   metrics=['accuracy'])
389     return model
390
391 resnet_model = build_resnet50_model()
392
393 history_resnet = resnet_model.fit(
394     X_train_resnet, y_train,
395     validation_data=(X_test_resnet, y_test),
396     epochs=10,
397     batch_size=32,
398     callbacks=[early_stop],
399     verbose=1

```

```

398 )
399
400
401 # ===== Cell 19 =====
402 resnet_preds = resnet_model.predict(X_test_resnet)
403 resnet_preds_classes = tf.argmax(resnet_preds, axis=1)
404
405 print("      Classification Report for ResNet50:")
406 print(classification_report(y_true, resnet_preds_classes))
407 print("      Accuracy:", accuracy_score(y_true, resnet_preds_classes))
408
409
410 # ===== Cell 20 =====
411 plot_history(history_resnet, model_name="ResNet50")
412
413
414 # ===== Cell 21 =====
415 cm = confusion_matrix(y_true, resnet_preds_classes)
416 disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=["Normal", "Diabetic"])
417 disp.plot(cmap='Oranges')
418 plt.title("Confusion Matrix - ResNet50")
419 plt.show()
420
421
422 # ===== Cell 22 =====
423 fpr, tpr, _ = roc_curve(y_test[:, 1], resnet_preds[:, 1])
424 roc_auc = auc(fpr, tpr)
425
426 plt.figure(figsize=(6, 6))
427 plt.plot(fpr, tpr, color='teal', lw=2, label='ROC curve (area = %0.4f)' % roc_auc)
428 plt.plot([0, 1], [0, 1], color='gray', lw=2, linestyle='--')
429 plt.title('ROC Curve - ResNet50')
430 plt.xlabel('False Positive Rate')
431 plt.ylabel('True Positive Rate')
432 plt.legend(loc="lower right")
433 plt.grid()
434 plt.show()
435
436
437 # ===== Cell 23 =====
438 from tensorflow.keras.applications import EfficientNetB0
439 from tensorflow.keras.applications.efficientnet import preprocess_input
440     as efficientnet_preprocess
441
442 # Prepare input
443 X_train_eff = expand_channels(X_train)
444 X_test_eff = expand_channels(X_test)
445
446 X_train_eff = efficientnet_preprocess(X_train_eff)
447 X_test_eff = efficientnet_preprocess(X_test_eff)
448
449 # Build EfficientNetB0 model with top layer fine-tuning
450 def build_efficientnet_model(input_shape=(224, 224, 3), num_classes=2):
451     base_model = EfficientNetB0(weights='imagenet', include_top=False,
452         input_shape=input_shape)

```

```

452     # Freeze all layers initially
453     for layer in base_model.layers:
454         layer.trainable = False
455
456     # Unfreeze top 20 layers for fine-tuning
457     for layer in base_model.layers[-20:]:
458         if not isinstance(layer, layers.BatchNormalization):
459             layer.trainable = True
460
461     x = layers.GlobalAveragePooling2D()(base_model.output)
462     x = layers.Dense(128, activation='relu')(x)
463     x = layers.Dropout(0.4)(x)
464     outputs = layers.Dense(num_classes, activation='softmax')(x)
465
466     model = Model(inputs=base_model.input, outputs=outputs)
467     model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=1e
468                 -4),
469                 loss='categorical_crossentropy',
470                 metrics=['accuracy'])
471
472     return model
473
474 efficientnet_model = build_efficientnet_model()
475
476 history_eff = efficientnet_model.fit(
477     X_train_eff, y_train,
478     validation_data=(X_test_eff, y_test),
479     epochs=10,
480     batch_size=32,
481     callbacks=[early_stop],
482     verbose=1
483 )
484
485 # ===== Cell 24 =====
486 eff_preds = efficientnet_model.predict(X_test_eff)
487 eff_preds_classes = tf.argmax(eff_preds, axis=1)
488
489 print("      Classification Report for EfficientNetB0:")
490 print(classification_report(y_true, eff_preds_classes))
491 print("      Accuracy:", accuracy_score(y_true, eff_preds_classes))
492
493 # ===== Cell 25 =====
494 plot_history(history_eff, model_name="EfficientNetB0")
495
496
497 # ===== Cell 26 =====
498 cm = confusion_matrix(y_true, eff_preds_classes)
499 disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=["Normal", "Diabetic"])
500 disp.plot(cmap='Greens')
501 plt.title("Confusion Matrix - EfficientNetB0")
502 plt.show()
503
504
505 # ===== Cell 27 =====
506 fpr, tpr, _ = roc_curve(y_test[:, 1], eff_preds[:, 1])
507 roc_auc = auc(fpr, tpr)

```

```

508 plt.figure(figsize=(6, 6))
509 plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area =
510     %0.4f)' % roc_auc)
511 plt.plot([0, 1], [0, 1], color='gray', lw=2, linestyle='--')
512 plt.title('ROC Curve - EfficientNetB0')
513 plt.xlabel('False Positive Rate')
514 plt.ylabel('True Positive Rate')
515 plt.legend(loc="lower right")
516 plt.grid()
517 plt.show()
518
519
520 # ===== Cell 28 =====
521 # Flatten the grayscale images (224x224) to vectors of shape (224*224, )
522 X_train_flat = X_train.reshape(X_train.shape[0], -1)
523 X_test_flat = X_test.reshape(X_test.shape[0], -1)
524
525 # Convert labels from one-hot to single integers
526 y_train_flat = np.argmax(y_train, axis=1)
527 y_test_flat = np.argmax(y_test, axis=1)
528
529
530 # ===== Cell 29 =====
531 from sklearn.decomposition import PCA
532
533 pca = PCA(n_components=50, random_state=42)
534 X_train_pca = pca.fit_transform(X_train_flat)
535 X_test_pca = pca.transform(X_test_flat)
536
537
538 # ===== Cell 30 =====
539 import umap
540
541 umap_reducer = umap.UMAP(n_components=50, random_state=42)
542 X_train_umap = umap_reducer.fit_transform(X_train_flat)
543 X_test_umap = umap_reducer.transform(X_test_flat)
544
545
546 # ===== Cell 31 =====
547 from sklearn.linear_model import LogisticRegression
548 from sklearn.svm import SVC
549 from sklearn.ensemble import RandomForestClassifier
550
551 def evaluate_classifiers(X_train_red, X_test_red, y_train_cls,
552     y_test_cls, reducer_name="PCA"):
553     print(f"\n      Evaluation with {reducer_name} Features:\n")
554
555     models = {
556         "Logistic Regression": LogisticRegression(max_iter=1000),
557         "SVM": SVC(kernel='rbf', probability=True),
558         "Random Forest": RandomForestClassifier(n_estimators=100,
559             random_state=42)
560     }
561
562     for name, model in models.items():
563         model.fit(X_train_red, y_train_cls)
564         preds = model.predict(X_test_red)

```

```

563     acc = accuracy_score(y_test_cls, preds)
564     print(f"      {name} Accuracy: {acc:.4f}")
565     print(classification_report(y_test_cls, preds))
566
567
568 # ===== Cell 32 =====
569 evaluate_classifiers(X_train_pca, X_test_pca, y_train_flat, y_test_flat,
570 , "PCA")
571
572 # ===== Cell 33 =====
573 evaluate_classifiers(X_train_umap, X_test_umap, y_train_flat,
574 y_test_flat, "UMAP")
575
576 # ===== Cell 34 =====
577 from sklearn.ensemble import VotingClassifier
578
579 # Instantiate individual models
580 log_clf = LogisticRegression(max_iter=1000, random_state=42)
581 svm_clf = SVC(probability=True, kernel='rbf', random_state=42)
582 rf_clf = RandomForestClassifier(n_estimators=100, random_state=42)
583
584 # Combine them into a hard-voting ensemble
585 voting_clf = VotingClassifier(
586     estimators=[
587         ('lr', log_clf),
588         ('svm', svm_clf),
589         ('rf', rf_clf)
590     ],
591     voting='hard'
592 )
593
594
595 # ===== Cell 35 =====
596 voting_clf.fit(X_train_umap, y_train_flat)
597
598
599 # ===== Cell 36 =====
600 from sklearn.metrics import accuracy_score, classification_report,
601     confusion_matrix, ConfusionMatrixDisplay
602
603 voting_preds = voting_clf.predict(X_test_umap)
604
605 print("      VotingClassifier Evaluation:")
606 print("      Accuracy:", accuracy_score(y_test_flat, voting_preds))
607 print(classification_report(y_test_flat, voting_preds))
608
609 # Confusion Matrix
610 cm = confusion_matrix(y_test_flat, voting_preds)
611 disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=["Normal", "Diabetic"])
612 disp.plot(cmap='Blues')
613 plt.title("Confusion Matrix - VotingClassifier (UMAP features)")
614 plt.grid(False)
615 plt.show()
616

```

```
617 | # ===== Cell 37 =====
```

Listing 1: Full Notebook Code with Cell Structure